



LF ENERGY

# OpenSynth

## Workshop: How to Train a Generative AI Model to Create Synthetic Smart Meter Data

Gus Chadney & Sheng Chai, Centre for Net Zero



Centre for Net Zero

Powered by Octopus Energy Group



# Today's Agenda

---

- 1 Training a generative model
  - 2 Evaluating Privacy
  - 3 Evaluating Fidelity and Utility
- 

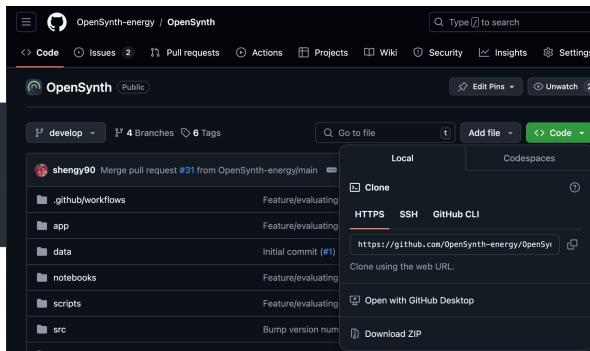


# Setting up your environment

**A** Clone the OpenSynth repository:  
<https://github.com/OpenSynth-energy/OpenSynth/>

**B** Install opensynth-energy with pip

```
1 pip install opensynth-energy
```



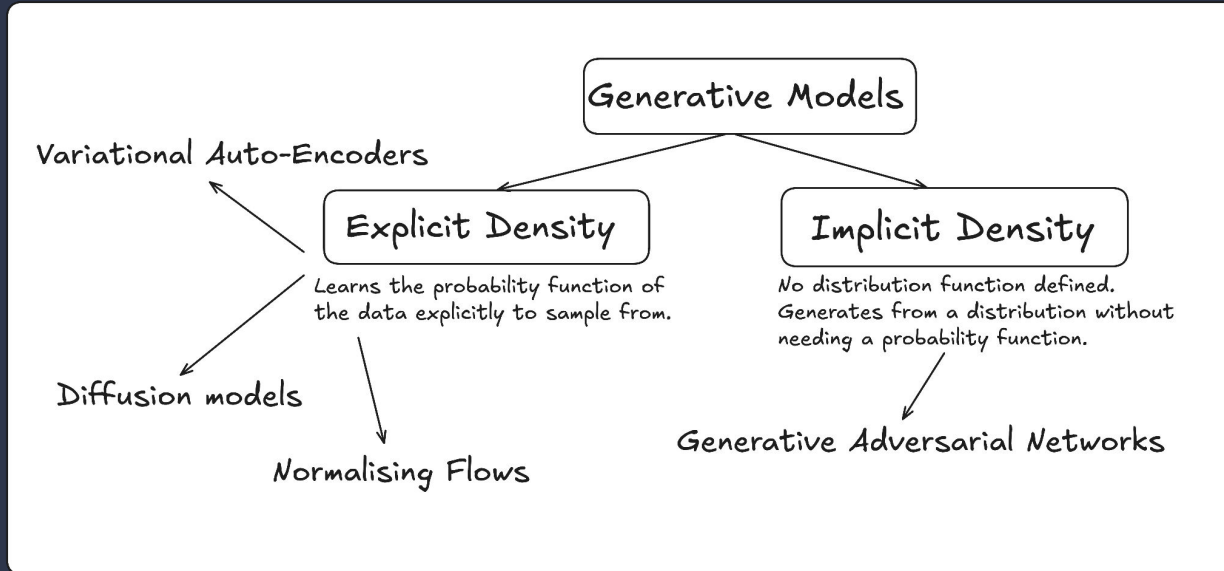
## Pre-requisites

1. Python version  $\geq 3.11$
2. Numpy version  $< 2.0.0$
3. Pytorch lightning  $\geq 2.3.1$
4. Pipenv for managing dependencies

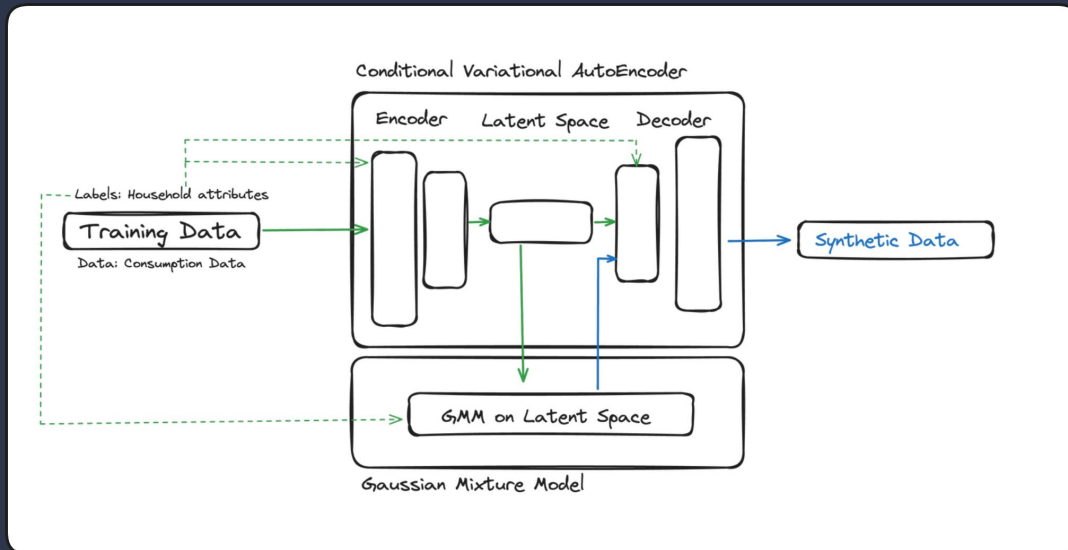
## Installing Pipenv + Repo

1. Pip install pipenv
2. Pipenv install torch" $\geq 2.3.1$ "
3. Pipenv sync --dev

# 1 Training a generative model



# Faraday - VAE + GMM



## Encoder

Maps raw data into a latent space.

## GMM

Learns distribution of the latent space.

## Decoder

During inference, samples from GMM and decode with decoder.



Faraday Paper

# Lab session: Training Faraday

## Low Carbon London Dataset

<https://shorturl.at/ahe3p>

## Notebooks

Github [repo/notebooks/faraday/faraday\\_tutorial.ipynb](https://github.com/centre-for-net-zero/repo/notebooks/faraday/faraday_tutorial.ipynb)

## If you need GPU

1. Switch to Google-Colab-Notebooks branch
2. Open [notebooks/faraday/Training\\_Generative\\_Model.ipynb](#) in Google Colab
3. Download trained model artefacts to your computer!

## If you have GPU

Check that Pytorch can access your GPU.

```
1 import torch
2
3 # If on windows (Nvidia GPUs)
4 torch.cuda.is_available()
5
6 # If on Mac machines
7 torch.backends.mps.is_available()
```



## Evaluating Privacy

**Privacy: Risk of private data leaking, i.e. original training data being recovered from generated samples.**

### Implicit Protection

During training, limit the model's ability to learn from data to prevent private data from being "memorised".

### Explicit Attacks

After training, prove that generated samples are indistinguishable **from unseen training data** to give "**plausible deniability**".





## Differential Privacy

Inclusion or exclusion of any single data record should not have “disproportionate” impact on the model’s output:

- Two datasets,  $D_1$  and  $D_2$  that differs by only 1 row
- $A$  is any algorithm, e.g. a ‘mean’  $\rightarrow A(D_1) = \text{Mean of } D_1$
- $S$  is all possible values of  $D_1$
- To be considered  **$(\epsilon, \delta)$ -differentially private**:

$$\Pr(A(D_1) \in S) \leq e^\epsilon \Pr(A(D_2) \in S) + \delta.$$

**Smaller  $\epsilon$  = more private.**

**$\delta$  typically set to  $< 1/N$  where  $N$  is dataset size.**







# Differential Private Stochastic Gradient Descent

- **ML Algorithm learns through SGD**

Stochastic gradient descent algorithm: updating model weights iteratively using randomly selected subsets (batches) of the training data to minimize a loss function.

- **DP-SGD modifies SGD to make process differentially private**

**DP-SGD** introduces **1) gradient clipping** to clip the largest gradient in a batch, typically set norm value to 1. and **2) injecting noise** to 'mask' the largest gradient in a batch and **3) (privacy-accounting)** keep track of how much noise was added.

In PyTorch, commonly implemented using the [Opacus](#) library.

**Smaller  $\epsilon$  = more private.**

**$\delta$  typically set to  $< 1/N$  where  $N$  is dataset size.**



'Defining Good' Paper





# Explicit Privacy Attacks

## Membership Inference Attack (MIA)

Distinguish between synthetic data vs unseen holdout data.

## Reconstruction attack

Force a generative model to produce output that resembles closely to original training data.

## Attribute Inference Attack

Reveal sensitive information from synthetic data.





# Membership Inference Attack

Consider the scenario. You have:

- **Model A:** predict whether someone has a criminal record
- **Model B:** predict whether someone belongs to the training data of Model A

If **Model B** predicts **Person A** belong to **Model B** with high confidence, and **Model A** predicts **Person A** has criminal records with high confidence, then we can have high confidence that **Person A** has criminal record → Leaking the privacy of **Person A**.

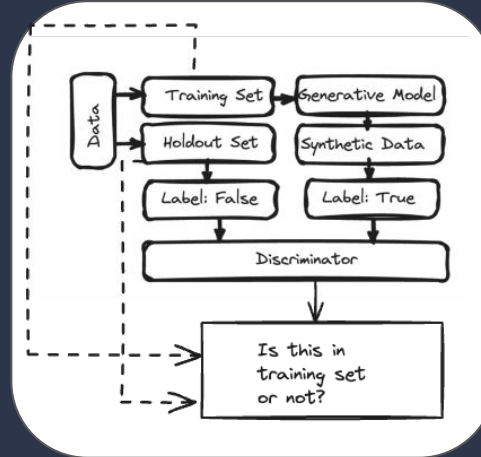
Distinguish between  
synthetic data vs  
unseen holdout data.




## 2

# Membership Inference Attack

1. **Split training data into Train and Holdout.** Train is used to train generative model. Holdout is used for privacy analysis.
2. **Generate a sample of synthetic data.** This data is labelled true. Holdout data is labelled false.
3. **Train a discriminator to distinguish between synthetic vs holdout data.**



**Holdout and Training data distribution however is too similar for traditional MIA to work!**

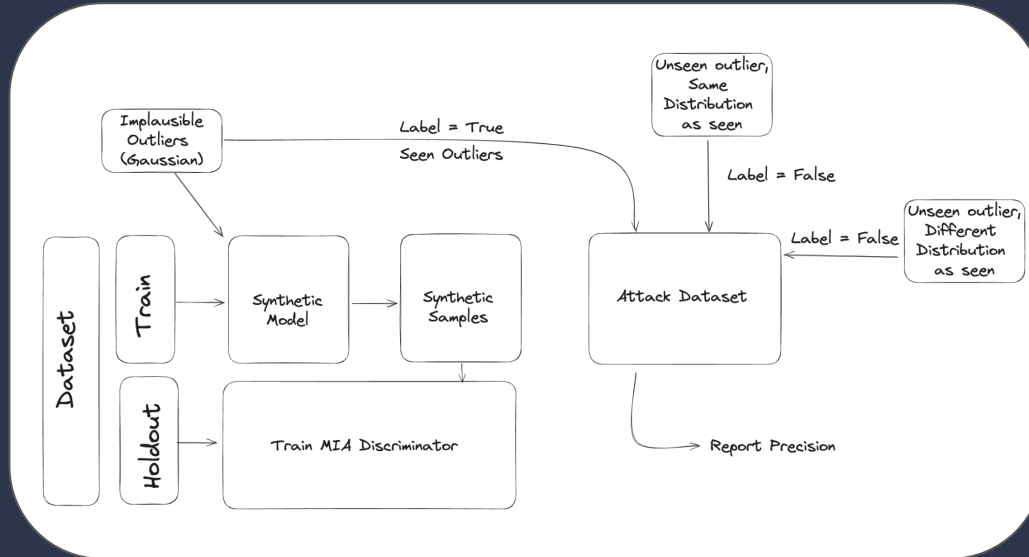


**Distinguish between synthetic data vs unseen holdout data.**



2

# Membership Inference Attack with Outliers




**Inject implausible outliers to training data for generative model.  
Then launch MIA directly on injected outliers.**

**Distinguish between  
synthetic data vs  
unseen holdout data.**



## 2

# Reconstruction Attack



One-tailed Two-sample KS Test:

Null Hypothesis: Synthetic Data is closer to Holdout than to training data.


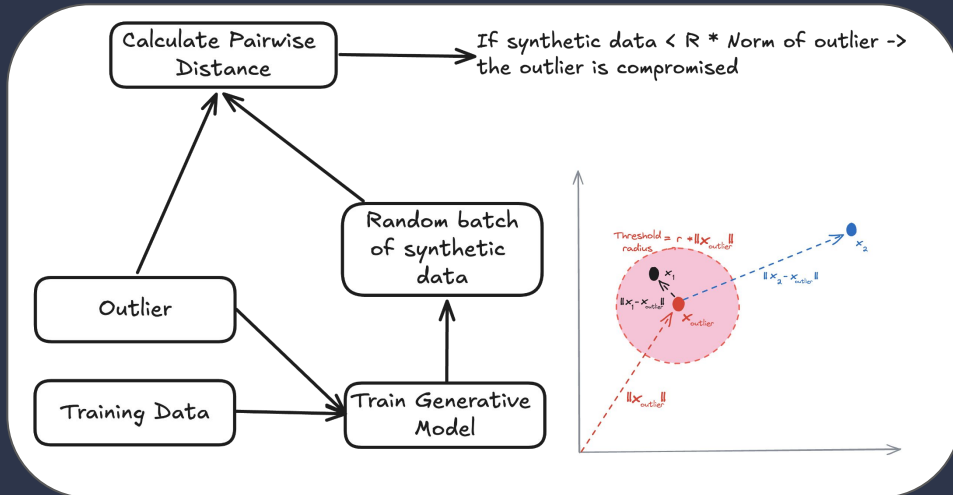
Alt Hypothesis: Synthetic Data is not closer to Holdout than to training data.

**Distance-based attack to prove that synthetic data is not more similar to training than to holdout data.**

**Force a generative model to produce output that resembles closely to original training data.**



## 2 Reconstruction Attack with Outliers




Stakeholders decide on a threshold radius.

Force a generative model to produce output that resembles closely to original training data.

Calculate % of outliers that were compromised under the decided threshold.

Ratio of 0.3 ~ Vector norm with 30 kWh: outlier with norm of  $30 \pm 9$  kWh would compromise it.



# Lab session: Evaluating Privacy

## Notebooks

Github [repo/notebooks/faraday/faraday\\_evaluation.ipynb](https://github.com/repo/notebooks/faraday/faraday_evaluation.ipynb)



‘Defining Good’ Paper

## You won't need GPU for this.

Repo won't work on Google Colab because of python compatibility.

Google colab uses 3.10.  
OpenSynth requires 3.11.

## If you have GPU

Check that Pytorch can access your GPU.

```
1 import torch
2
3 # If on windows (Nvidia GPUs)
4 torch.cuda.is_available()
5
6 # If on Mac machines
7 torch.backends.mps.is_available()
```



## 3 Evaluating Fidelity and Utility

### Fidelity

How similar synthetic data is statistically to original training data. Evaluated with statistical metrics, e.g. mean, quantiles, autocorrelation etc.

### Utility

Would we still get similar results if we used synthetic data in place of real data for downstream applications?





## Proposed Fidelity Metrics

**1. Standard statistical and distance measures:** Calculate the mean and 95th of each load profile. Compare the distributions of mean/ quantile values between Real and Synthetic data.

**2. Distribution of autocorrelation function (ACF) :** Calculate the ACF coefficients of each load profile. Compare the distributions of ACF coefficients between Real and Synthetic data.

**3. Distribution of magnitude and timing of peaks :** Mask each load profile with 0 and preserving only the 'Top K' values. Compare the distribution of masked dataset between Real and Synthetic data.

**4. Distribution of clusters:** Fit N arbitrary clusters using Gaussian Mixture Model on real data to predict clusters, and obtained predicted clusters of synthetic data. Compare distribution of clusters between Real and Synthetic Data.

**Smart meter data is hierarchical. Repeat 1-3 on a cluster level using clusters from 4.**



## 3

# Proposed Utility Metrics

Train on Synthetic, Test on Real Framework. For a downstream task:

1. Train **Model A** on **Real data**
2. Train **Model B** on **Synthetic Data**
3. Compare performances of **Model A** and **Model B** on unseen real data.

## Classification task

Predict season using smart meter data.

## Forecasting task

Predict next-time step consumption using historical data.



# Lab session: Evaluating Fidelity

 **Implement your own fidelity and utility metrics!**

 **Prepared Dataset**

<https://drive.google.com/drive/u/0/folders/1RJV-OFgelPwWVYkTK6YPQ2TSwu4LVreN>



‘Defining Good’ Paper

## If you need GPU

1. Create a new google colab notebook
2. Download the prepared dataset for fidelity/ utility analysis, upload it manually to Google Colab and have a go!

## If you have GPU

Check that Pytorch can access your GPU.

```
1 import torch
2
3 # If on windows (Nvidia GPUs)
4 torch.cuda.is_available()
5
6 # If on Mac machines
7 torch.backends.mps.is_available()
```